

Making things go right
Fixing things that go wrong

Making things go right

- Defensive programming : *aim for the pit of success*
 - Creating abstractions : *manage and protect state*
 - Adding assertions : *what are you assuming is true?*
 - Creating test-benches : *what is the correct result?*
 - Incremental changes : *don't break existing functions*
- Other techniques
 - Compile with warnings turned on: **fix** compiler warnings
 - Specify your functions in the abstract
 - Develop high-level algorithms
 - Separate complex processes into smaller functions

Specifying functions : sorting

"Sort a range of integers"

```
void sort(int *begin, int *end);  
void sort(vector<int> &begin);  
vector<int> sort(const vector<int> &begin);
```

Specifying functions : sorting

"Sort a range of integers"

```
void sort(int *begin, int *end);  
void sort(vector<int> &begin);  
vector<int> sort(const vector<int> &begin);
```

Specifying functions : sorting

"Sort a range of integers"

```
bool is_sorted(const int *begin, const int *end);
```

```
// Pre:  begin <= end
```

```
// Post: is_sorted(begin',end')
```

```
void sort(int *begin, int *end)
```

```
{
```

```
    assert(begin <= end);
```

```
    // Something
```

```
    assert(is_sorted(begin, end));
```

```
}
```

Specifying functions : sorting

"Sort a range of integers"

```
bool is_sorted(const int *begin, const int *end);
```

```
// Pre:  begin <= end
```

```
// Post: is_sorted(begin',end')
```

```
void sort(int *begin, int *end)
```

```
{
```

```
    assert(begin <= end); // Pre-condition
```

```
    // Something
```

```
    assert(is_sorted(begin, end)); // Post-condition
```

```
}
```

Specifying functions : induction

```
// Pre:  is_sorted(begin,end-1)
// Post: is_sorted(begin',end') && end[-1] in (begin',end')
void insert(int *begin, int *end);

void sort(int *begin, int *end)
{
    assert(begin <= end); // Pre-condition

    if(begin==end){
        // An empty list is sorted
    }else{
        sort(begin, end-1); // Sort range [begin,end-1)
        insert(begin, end); // Produce sorted [begin,end)
    }

    assert(is_sorted(begin, end)); // Post-condition
}
```

Specifying functions : induction

```
// Pre:  is_sorted(begin,end-1)
// Post: is_sorted(begin',end') && end[-1] in (begin',end')
void insert(int *begin, int *end);

void sort(int *begin, int *end)
{
    assert(begin <= end); // Pre-condition

    if(begin==end){
        // An empty list is sorted
    }else{
        sort(begin, end-1); // Sort range [begin,end-1)
        insert(begin, end); // Produce sorted [begin,end)
    }

    assert(is_sorted(begin, end)); // Post-condition
}
```


Specifying functions : induction

```
// Insert end[-1] into sorted range[begin,end-1)
// Pre:  is_sorted(begin,end-1)
// Post: is_sorted(begin',end') && end[-1] in (begin',end')
void insert(int *begin, int *end)
{
    assert( is_sorted(begin,end-1) );

    int i=begin-end-1;
    while( 0 < i){
        if( begin[i-1] > begin[i] ){
            swap( begin[i-1], begin[i] );
        }
        i=i-1;
    }

    assert( is_sorted(begin,end) );
}
```

Is this code correct?

Specifying functions : testing

```
// Insert end[-1] into sorted range[begin,end-1)
// Pre:  is_sorted(begin,end-1)
// Post: is_sorted(begin',end') && end[-1] in (begin',end')
void insert(int *begin, int *end)
{
    assert( is_sorted(begin,end-1) );

    int i=begin-end-1;
    while( 0 < i){
        if( begin[i-1] > begin[i] ){
            swap( begin[i-1], begin[i] );
        }
        i=i-1;
    }

    assert( is_sorted(begin,end) );
}
```

Is this code correct?

What inputs to test?

- Designing test cases is a bit of an art
 - Think of special cases
 - Think of edge cases
 - Include stress tests
- Test cases for sorting a vector
 - Empty vector : []
 - Single element vector : [a]
 - Two element vector
 - Ordered : [a,b], $a < b$
 - Un-ordered: [a,b], $a > b$
 - Equal [a,b], $a == b$
 - Three element vector
 - All combinations of ordered, un-ordered, and equal?
 - Random vectors

Pre and post-conditions in practise

- Pre- and post-conditions are extremely useful
 - Documentation: what is the function supposed to do?
 - Testing: what are valid inputs and outputs?
 - Debugging: when does the function fail?
- Asserting pre/post-conditions *may* be useful
 - Are the conditions cheap or expensive?
 - Is checking the condition slower than the function?
 - Assert the cheap stuff: is this pointer non-null?
 - Comment out the expensive stuff: is this array sorted?
 - But leave it in as documentation and for future debugging

Fixing things that go wrong

- Defensive programming
 - Enforce abstractions : *check internal state is valid*
 - Adding more assertions : *how early does problem occur?*
 - Extend test-benches : *what is the simplest failing case?*
- Other techniques
 - Logging via stderr
 - Debugging
 - Exceptions

"My program doesn't work"

- Programs can "not work" in many different ways
- In what way is it not working?
 - **Compilation:** the source code is malformed
 - **Linking:** definitions are missing or repeated
 - **Crashing:** program does not complete execution
 - **Correctness:** program executes, but output is wrong

Running the wrong program

*"I keep changing the source code,
but the program still fails in the same way"*

1. Is compilation actually completing successfully?
2. Are you running the program/function you just compiled?
3. Is the code you are editing actually getting compiled?
4. Are you in the right directory?

Suggestion: make the program print something unique

```
int main(int argc, char *argv[])
{
    cerr << "Debug session 1456" << endl;

    assert(argc > 1);
    int x = atoi(argv[1]);
}
```

Dealing with compiler errors

- C++ produces incredibly long error messages
 - They are notoriously difficult to comprehend
 - A single error can result in hundreds of lines of error
 - Templates mean there is an explosion of `::` and `<>`
- **Tip 1** : fix the ***first*** error message first
 - *Ignore the last error*; it could be caused by the first one
 - Scroll up to the very first error message
- **Tip 2** : ***read*** the error messages (with eyes+brain)
 - The compiler is trying to give you all the information it can
 - Look at all the source locations it identifies very carefully
- **Tip 3** : google the error messages
 - Only do this after trying to solve the problem yourself
 - Actually read the explanation for what went wrong
 - Beware of "cargo-cult" programming

Compiling versus linking

- A program is compiled in two main phases:
 1. **Compiling**: turn source files (.cpp) into object files (.o)
 2. **Linking**: link object files (.o) together into one program
- Object file contains the "machine-code" for a source file
 - Meta-data about the declarations and definitions
 - CPU-specific binary instructions for the defined functions
 - Source file -> object file : easy; performed by the compiler
 - Object file -> source file : hard; no general solution
- Compilers will both compile and link by default
 - Compile and link: `g++ fileA.cpp fileB.o -o prog`
 - Compile to object file only: `g++ -c fileA.cpp -o fileA.o`
 - Link in existing object file: `g++ fileA.o fileB.cpp -o prog`

Dealing with linker errors

- Linker errors are generally easier to deal with
 - Linking is just about matching decls. and defs.
- 1. "undefined reference": declared but not defined
 - Something is declared somewhere and used
 - That thing is never defined in any source or object file
 1. You forgot to add a source/object file while compiling; or
 2. You forgot to add a definition; or
 3. The definition does not match the declaration
- 2. "multiple definition" : something is defined twice
 - You have compiled in the same source file twice
 - Multiple headers provide a definition of a function
 - Split function up: declaration in header; definition in one source

Crashes: "program go boom"

- A "crash" is a loose term for lots of things
 - The program did not exit in a controlled way; and/or
 - Some sort of exception or run-time error happened
 - It may not be obvious that the program has crashed
- We have four broad classes of crash
 1. **Stack overflow**: a function recurses too many times
 2. **Memory corruption**: reading or writing via invalid pointers
 3. **Undefined behavior**: stepping outside the rules of C++
 4. **Unhandled exceptions**: C++ mechanism to indicate errors

Crashes : using debuggers

- A debugger allows you to explore program execution
 - *Where did the program crash?*
 - *What functions are on the stack?*
 - *What are the values of local variables?*
- There are many different debuggers available
 - Command-line: gdb (unix/linux), lldb (OSX)
 - GUI: Visual Studio, XCode, CLion, ...
 - Debuggers are very powerful and can be complicated
- *Recommendation*: use command-line gdb
 - Always available in Linux
 - Can answer the key "where" and "what" questions

Stack overflow : the problem

```
int f(int n)
{
    if(n==0)
        return 1;
    return f(n-1)+f(n-2);
}
```

```
int main()
{
    return f(10);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ g++ fib.cpp -o fib
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ ./fib
```

```
Segmentation fault (core dumped)
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$
```

Stack overflow : adding debug info

```
int f(int n)
{
    if(n==0)
        return 1;
    return f(n-1)+f(n-2);
}
```

```
int main()
{
    return f(10);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ g++ -g fib.cpp -o fib
```

```
dt10@LAPTOP-0DEHDEQ0:~
$
```

Stack overflow : running in gdb

```
int f(int n)
{
    if(n==0)
        return 1;
    return f(n-1)+f(n-2);
}

int main()
{
    return f(10);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ g++ -g fib.cpp -o fib
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ gdb ./fib
```

```
GNU gdb (Ubuntu 8.1-0ubuntu3)
```

```
8.1.0.20180409-git
```

```
...
```

```
Reading symbols from ./fib...done.
```

```
(gdb)
```

Stack overflow : running in gdb

```
int f(int n)
{
    if(n==0)
        return 1;
    return f(n-1)+f(n-2);
}

int main()
{
    return f(10);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ g++ -g fib.cpp -o fib
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ gdb ./fib
```

```
GNU gdb (Ubuntu 8.1-0ubuntu3)
```

```
8.1.0.20180409-git
```

```
...
```

```
Reading symbols from ./fib...done.
```

```
(gdb) run
```


Stack overflow : running in gdb

```
int f(int n)
{
    if(n==0)
        return 1;
    return f(n-1)+f(n-2);
}

int main()
{
    return f(10);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ g++ -g fib.cpp -o fib
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ gdb ./fib
```

```
GNU gdb (Ubuntu 8.1-0ubuntu3)
```

```
8.1.0.20180409-git
```

```
...
```

```
Reading symbols from ./fib...done.
```

```
(gdb) run
```

```
Starting program: ~/fib
```

```
Program received signal SIGSEGV,  
Segmentation fault.
```

```
0x000000008000603 in f (...) at
```

```
fib.cpp:2
```

```
2      {
```

Stack overflow : viewing the stack

```
int f(int n)
{
    if(n==0)
        return 1;
    return f(n-1)+f(n-2);
}
```

```
int main()
{
    return f(10);
}
```

(gdb) run

Starting program: ~/fib

Program received signal SIGSEGV,
Segmentation fault.

0x000000008000603 in f (...) at
fib.cpp:2

2 {

(gdb) bt

#0 0x000000008000603 in f (n=<error>)
at fib.cpp:2

#1 0x000000008000620 in f (n=-174672)
at fib.cpp:6

#2 0x000000008000620 in f (n=-174671)
at fib.cpp:6

#3 0x000000008000620 in f (n=-174670)
at fib.cpp:6

#4 0x000000008000620 in f (n=-174669)

Corruption: the problem

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ g++ tmp.cpp
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ gdb 0 Hello
H
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ gdb 1 Hello
e
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ gdb 100000 Hello
Segmentation fault (core dumped)
```

```
dt10@LAPTOP-0DEHDEQ0:~
$
```

Corruption: adding debug info

```
char f(int off, char *p)
{
    return p[off];
}
```

```
int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ g++ -g tmp.cpp
```

```
dt10@LAPTOP-0DEHDEQ0:~
$
```

Corruption: running in gdb

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ g++ -g tmp.cpp
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ gdb --args ./a.out 10000 Hello
```

Corruption: running in gdb

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ g++ -g tmp.cpp
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ gdb --args ./a.out 10000 Hello
```

```
GNU gdb (Ubuntu 8.1-0ubuntu3)
```

```
8.1.0.20180409-git
```

```
(gdb)
```

Corruption: running in gdb

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ g++ -g tmp.cpp
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ gdb --args ./a.out 10000 Hello
```

```
GNU gdb (Ubuntu 8.1-0ubuntu3)
```

```
8.1.0.20180409-git
```

```
(gdb) run
```

```
Starting program: ~/a.out 10000 Hello
```

```
Program received signal SIGSEGV,  
Segmentation fault.
```

```
0x000000008000862 in f (off=10000,  
p=0x7fffffffefee451 "Hello") at tmp.cpp:7  
7          return p[off];
```

```
(gdb)
```

Corruption: running in gdb

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ g++ -g tmp.cpp
```

```
dt10@LAPTOP-0DEHDEQ0:~
```

```
$ gdb --args ./a.out 10000 Hello
```

```
GNU gdb (Ubuntu 8.1-0ubuntu3)
```

```
8.1.0.20180409-git
```

```
(gdb) run
```

```
Starting program: ~/a.out 10000 Hello
```

```
Program received signal SIGSEGV,  
Segmentation fault.
```

```
0x000000008000862 in f (off=10000,  
p=0x7fffffffefee451 "Hello") at tmp.cpp:7  
7          return p[off];
```

```
(gdb)
```


Corruption: running in gdb

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ g++ -g tmp.cpp
```

```
dt10@LAPTOP-0DEHDEQ0:~
$ gdb --args ./a.out 10000 Hello
GNU gdb (Ubuntu 8.1-0ubuntu3)
8.1.0.20180409-git
```

```
(gdb) run
Starting program: ~/a.out 10000 Hello
```

```
Program received signal SIGSEGV,
Segmentation fault.
0x000000008000862 in f (off=10000,
p=0x7fffffffefee451 "Hello") at tmp.cpp:7
7          return p[off];
(gdb)
```

Corruption: viewing the stack

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

(gdb) run

Starting program: ~/a.out 10000 Hello

Program received signal SIGSEGV,
Segmentation fault.

0x0000000008000862 in f (off=10000,
p=0x7fffffffefee451 "Hello") at tmp.cpp:7
7 return p[off];

(gdb)

Corruption: viewing the stack

```
char f(int off, char *p)
{
    return p[off];
}

int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

(gdb) run

Starting program: ~/a.out 10000 Hello

Program received signal SIGSEGV,
Segmentation fault.

0x0000000008000862 inf (off=10000,
p=0x7fffffee451 "Hello") at tmp.cpp:7
7 return p[off];

(gdb) bt

#0 0x0000000008000862 in
f (off=10000, p=0x7fffffee451 "Hello")
at tmp.cpp:7

#1 0x00000000080008a4 in
main (argc=3, argv=0x7ffffee1c8)
at tmp.cpp:13

Corruption: viewing variables

```
char f(int off, char *p)
{
    return p[off];
}
```


```
int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

(gdb) bt

#0 0x0000000008000862 in
f (off=10000, p=0x7ffffee451 "Hello")
at tmp.cpp:7

#1 0x00000000080008a4 in
main (argc=3, argv=0x7ffffee1c8)
at tmp.cpp:13

(gdb)



Corruption: viewing variables

```
char f(int off, char *p)
{
    return p[off];
}
```

```
int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

(gdb) bt

#0 0x0000000008000862 in
f (off=10000, p=0x7ffffee451 "Hello")
at tmp.cpp:7

#1 0x00000000080008a4 in
main (argc=3, argv=0x7ffffee1c8)
at tmp.cpp:13

(gdb) print off
\$1 = 10000

(gdb) print p
\$2 = 0x7ffffee451 "Hello"

(gdb)

Corruption: moving up the stack

```
char f(int off, char *p)
{
    return p[off];
}
```

```
int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

(gdb) bt

```
#0  0x0000000008000862 in
f (off=10000, p=0x7ffffee451 "Hello")
at tmp.cpp:7
```

```
#1  0x00000000080008a4 in
main (argc=3, argv=0x7ffffee1c8)
at tmp.cpp:13
```

(gdb) up

```
#1  0x00000000080008a4 in
main (argc=3, argv=0x7ffffee1c8)
at tmp.cpp:13
```

```
13      cout << f(off, argv[2]);
```

(gdb)

Corruption: viewing variables

```
char f(int off, char *p)
{
    return p[off];
}
```

```
int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
(gdb) up
#1  0x00000000080008a4 in
main (argc=3, argv=0x7ffffee1c8)
    at tmp.cpp:13
13      cout << f(off, argv[2]);
```

(gdb)



Corruption: viewing variables

```
char f(int off, char *p)
{
    return p[off];
}
```

```
int main(int argc, char **argv)
{
    int off=atoi(argv[1]);
    cout << f(off, argv[2]);
}
```

```
(gdb) up
#1  0x00000000080008a4 in
main (argc=3, argv=0x7ffffee1c8)
at tmp.cpp:13
13      cout << f(off, argv[2]);
```

```
(gdb) print off
print off
$3 = 10000
```

```
(gdb) print argv[2]
$4 = 0x7ffffee451 "Hello"
```

```
(gdb)
```


Using a debugger

- Gdb is good for working out "where" and "what"
 - `run` : where did the program fail?
 - `bt` : (back-trace) what functions were on the stack?
 - `print` : what are the current values of variables?
 - `up` : moving to an outer function call
- Debuggers will not give you the "why"
 - It's up to you to work out what the actual root cause is
- Debuggers are usually a last resort
 - You've already thought about the code
 - You've already added assertions to check assumptions
 - You've already tried adding logging with `cerr`

Undefined behavior

- C++ allows you to shoot yourself in the foot
 - Not all behavior is specified
 - Some behavior is explicitly undefined
 - This allows compilers to create faster programs
- Some undefined behaviour *may* crash
 - Dereferencing a null pointer
 - Accessing memory after it is freed
 - Not returning a value from a function
- Some undefined behaviour *may not* crash
 - Dereferencing a null pointer
 - Accessing memory after it is freed
 - Not returning a value from a function

Undefined behavior : avoidance

- The compiler can *try* to help you
 - Compile with warnings turned on : `g++ -W -Wall`
 - Use run-time sanitisers: `g++ -fsanitize=undefined`
- There is no guaranteed method here
 - Create multiple test inputs
 - Examine every compiler warning
 - Try to diagnose and understand errors; don't use patch over